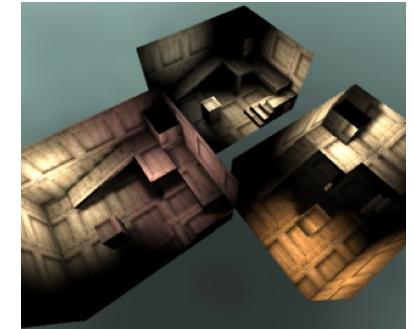


# Applications of Texture Mapping

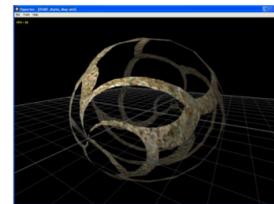
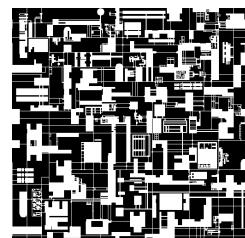
Prof.Vladlen Koltun  
Computer Science Department  
Stanford University

## Light maps

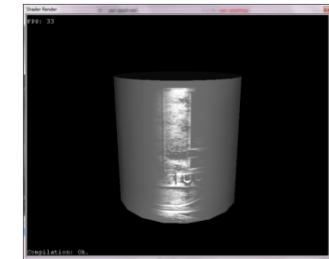
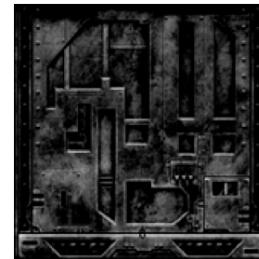


[http://www.Irrlicht3d.org](http://www Irrlicht3d.org)

## Opacity mapping

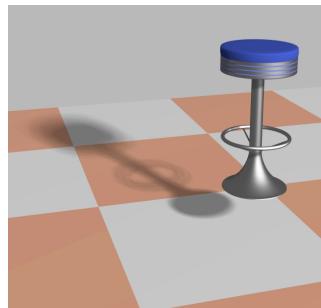
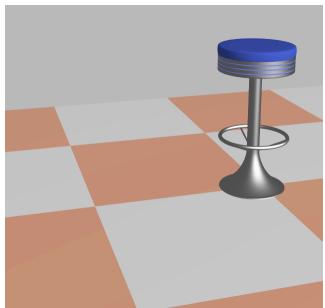


## Specular mapping



<http://lectrablog.lectralizard.com/page/2>

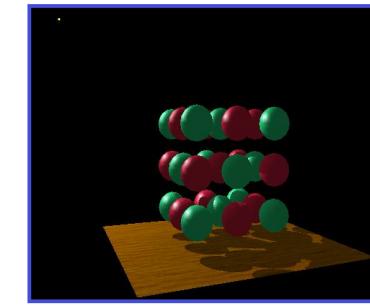
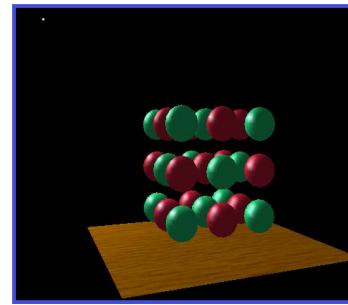
## Shadows



- Valuable cue of spatial relationships
- Increases realism

Akenine-Moeller and Haines

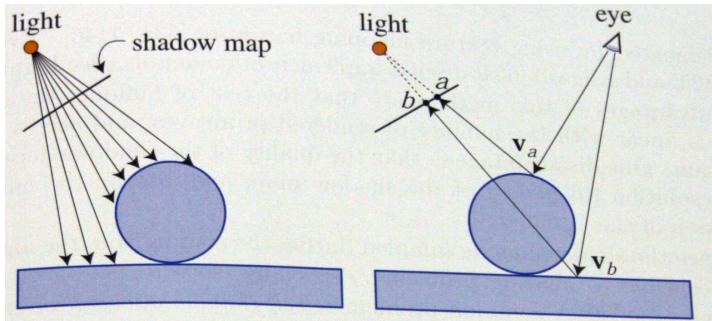
## Shadows



- Valuable cue of spatial relationships
- Increases realism

Mark J. Kilgard

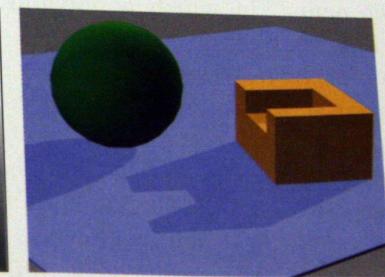
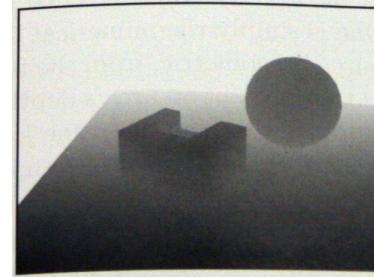
## Shadow mapping



- First pass: render the scene from the viewpoint of the light, store depth buffer as texture (shadow map)
- Second pass: project vertices into shadow map and compare depth values

Akenine-Moeller et al., Real-Time Rendering

## Shadow mapping



- First pass details: can disable all rendering features that do not affect depth map.
- Second pass details: For each fragment, use the light's modelview and projection transforms to obtain (u,v) coordinates in the shadow map and the depth w of the vertex.
- Compare w with value w' stored in (u,v) in the shadow map. If  $w \leq w'$ , perform lighting calculations with this light. Otherwise, do not.

Akenine-Moeller et al., Real-Time Rendering

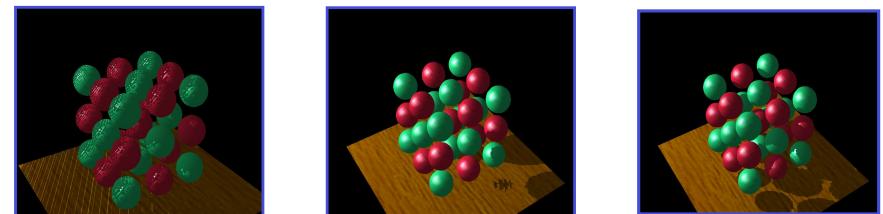
## Bias



- Numerical imprecision leads to self-shadowing
- Solution: add a bias  $\varepsilon$ . Change comparison from  $w \leq w'$  to  $w \leq w' + \varepsilon$
- Can use `glPolygonOffset`

Akenine-Moeller et al., Real-Time Rendering

## Setting the bias

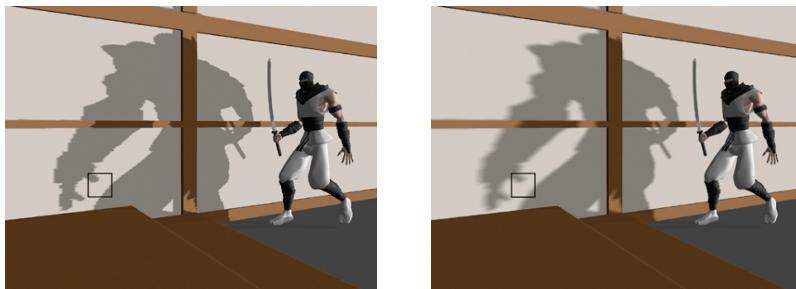


Too little                      Too much                      Just right

- Numerical imprecision leads to self-shadowing
- Solution: add a bias  $\varepsilon$ . Change comparison from  $w \leq w'$  to  $w \leq w' + \varepsilon$
- Can use `glPolygonOffset`

Mark J. Kilgard

## Shadow map aliasing



Unfiltered

Filtered

- Insufficient shadow map resolution leads to blocky shadows
- No easy solution. Should not filter depth values: leads to errors at object boundaries
- Percentage-closer filtering: filter comparison results

Bunnell and Pellacini

## Other issues

- Additional rendering pass for each shadow-casting light
- Setting the “field of view” of the light. Can use spotlights, or a cube map (six shadow maps) for a point light.
- For directional lights, use orthographic projection

## Reflection mapping



- Render the scene from a single point inside the reflective object. Store rendered images as textures.
- Map textures onto object. Determine texture coordinates by reflecting view ray about the normal.

Terminator 2

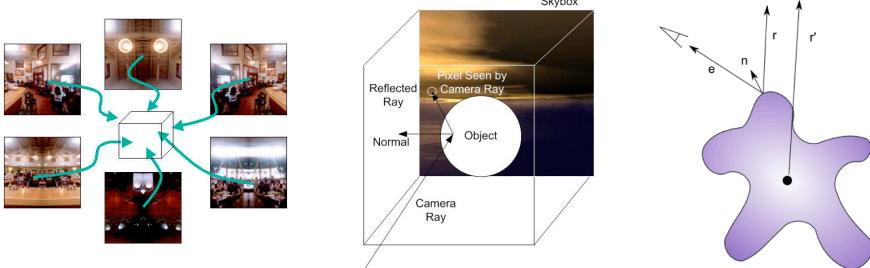
## Cube mapping



- Render the scene six times, through six faces of a cube, with 90-degree field-of-view for each image.
- Store images in six textures, which represent an omni-directional view of the environment

Greene, 1986

## Cube mapping



- To compute texture coordinates, reflect the view vector  $\mathbf{v}$  about the normal  $\mathbf{n}$ :
- $$\mathbf{r} = 2(\mathbf{v} \cdot \mathbf{n})\mathbf{n} - \mathbf{v}$$
- The highest (in absolute value) coordinate of  $\mathbf{r}$  identifies which of the six maps we need. The texture coordinates in this map are obtained by normalizing the other two coordinates of  $\mathbf{r}$ .

[http://developer.nvidia.com/object/cube\\_map\\_ogl\\_tutorial.html](http://developer.nvidia.com/object/cube_map_ogl_tutorial.html); TopherTG (Wikipedia)

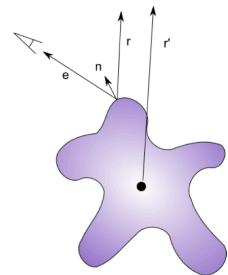
## Sphere mapping



- Cube maps require maintaining six texture in memory
- Sphere mapping uses a single viewpoint-specific environment map, updated every frame
- Map depicts a perfectly reflective sphere viewed orthographically

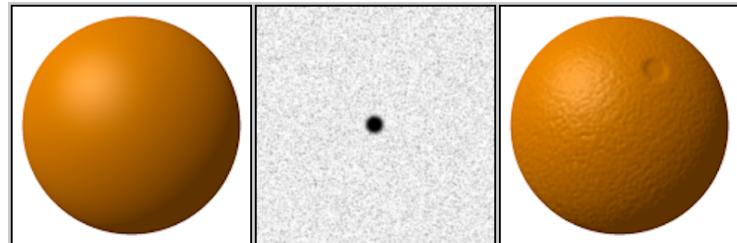
Greene, 1986

## Reflection mapping limitations



- Self-reflections not supported. A concave object will not reflect parts of itself.
- Environment map only correct for the point from which it was rendered. Good approximation for distant reflected objects, but can lead to substantial artifacts in general.

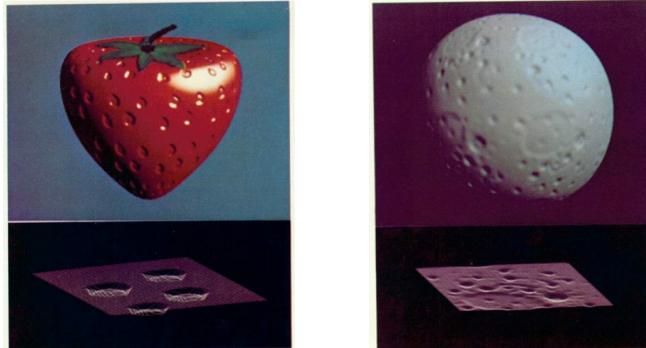
## Bump mapping



- Simulates roughness (“bumpiness”) of a surface without adding geometry
- Uses a two-dimensional height field (bump map) to perturb the normal during per-fragment shading calculations
- Limitation: silhouette is unaffected

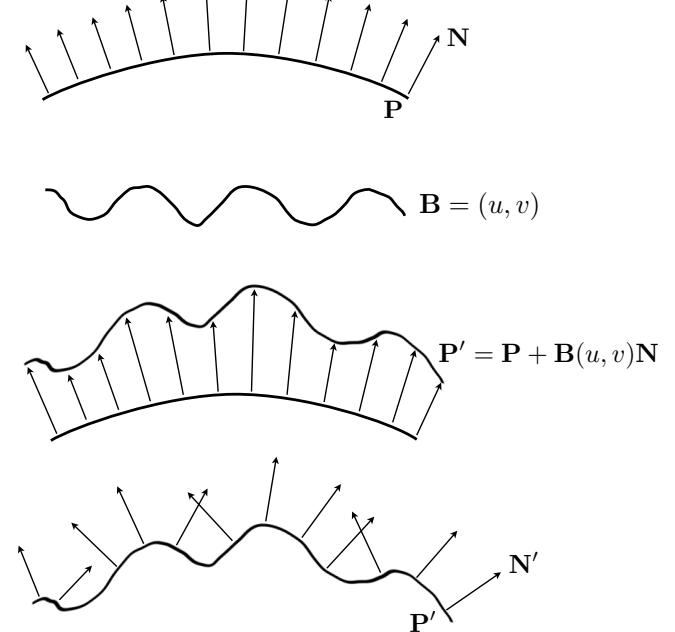
GDallimore (Wikimedia Commons)

## Bump mapping

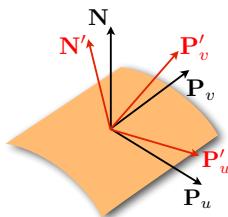


- Simulates roughness (“bumpiness”) of a surface without adding geometry
- Uses a two-dimensional height field (bump map) to perturb the normal during per-fragment shading calculations
- Limitation: silhouette is unaffected

Blinn, SIGGRAPH 1978



## Bump mapping derivation



$$\mathbf{N} = \mathbf{P}_u \times \mathbf{P}_v$$

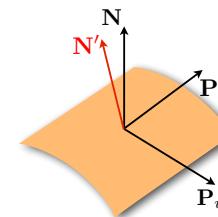
$$\mathbf{N}' = \mathbf{P}'_u \times \mathbf{P}'_v$$

The normals are always appropriately normalized, but we omit that for simplicity

$$\begin{aligned}\mathbf{P}'_u &= \frac{\partial \mathbf{P}'}{\partial u} & \mathbf{P}'_v &= \frac{\partial \mathbf{P}'}{\partial v} \\ &= \frac{\partial}{\partial u} (\mathbf{P} + \mathbf{B}(u, v)\mathbf{N}) & &= \frac{\partial}{\partial v} (\mathbf{P} + \mathbf{B}(u, v)\mathbf{N}) \\ &= \mathbf{P}_u + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} + \mathbf{B}(u, v) \frac{\partial \mathbf{N}}{\partial u} & &= \mathbf{P}_v + \frac{\partial \mathbf{B}}{\partial v} \mathbf{N} + \mathbf{B}(u, v) \frac{\partial \mathbf{N}}{\partial v} \\ &\approx \mathbf{P}_u + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} & &\approx \mathbf{P}_v + \frac{\partial \mathbf{B}}{\partial v} \mathbf{N}\end{aligned}$$

These are not really partial derivatives, but we follow Blinn's original notation

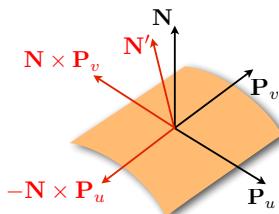
## Bump mapping derivation



$$\begin{aligned}\mathbf{N}' &= \mathbf{P}'_u \times \mathbf{P}'_v \\ &= \left( \mathbf{P}_u + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} \right) \times \left( \mathbf{P}_v + \frac{\partial \mathbf{B}}{\partial v} \mathbf{N} \right) \\ &= \mathbf{N} + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} \times \mathbf{P}_v + \frac{\partial \mathbf{B}}{\partial v} \mathbf{P}_u \times \mathbf{N} \\ &= \mathbf{N} + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} \times \mathbf{P}_v - \frac{\partial \mathbf{B}}{\partial v} \mathbf{N} \times \mathbf{P}_u\end{aligned}$$

The "partial derivatives" are just differences between adjacent pixel values in the bump map

## Bump mapping derivation

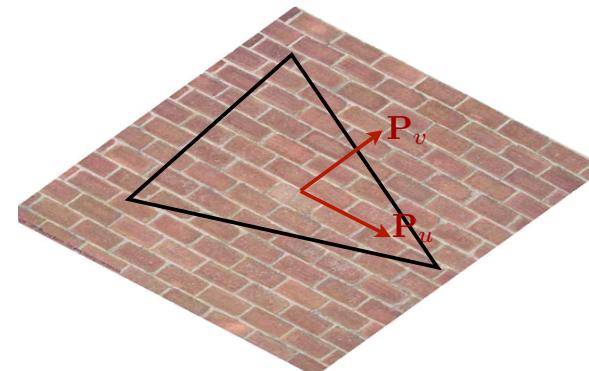


$$\begin{aligned}\mathbf{N}' &= \mathbf{P}'_u \times \mathbf{P}'_v \\ &= \left( \mathbf{P}_u + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} \right) \times \left( \mathbf{P}_v + \frac{\partial \mathbf{B}}{\partial v} \mathbf{N} \right) \\ &= \mathbf{N} + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} \times \mathbf{P}_v + \frac{\partial \mathbf{B}}{\partial v} \mathbf{P}_u \times \mathbf{N} \\ &= \mathbf{N} + \frac{\partial \mathbf{B}}{\partial u} \mathbf{N} \times \mathbf{P}_v - \frac{\partial \mathbf{B}}{\partial v} \mathbf{N} \times \mathbf{P}_u\end{aligned}$$

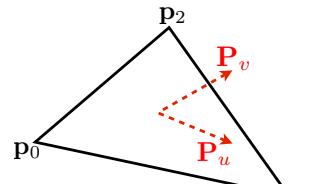
The "partial derivatives" are just differences between adjacent pixel values in the bump map

## Computing tangent space basis vectors

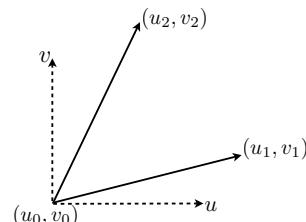
We now need to compute the vectors  $\mathbf{P}_u$  and  $\mathbf{P}_v$ . These are the directions on the surface that correspond to zero change in the  $v$  parameter and the  $u$  parameter, respectively.



## Computing tangent space basis vectors

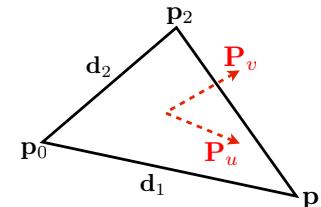


object space

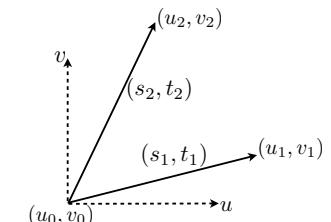


texture space

## Computing tangent space basis vectors



object space



texture space

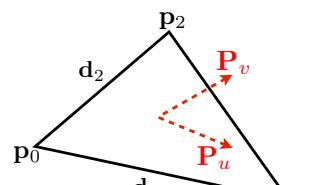
$$\mathbf{d}_1 = \mathbf{p}_1 - \mathbf{p}_0$$

$$\mathbf{d}_2 = \mathbf{p}_2 - \mathbf{p}_0$$

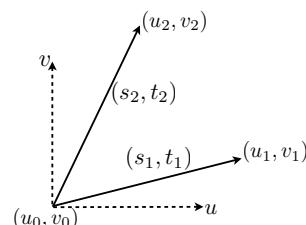
$$(s_1, t_1) = (u_1 - u_0, v_1 - v_0)$$

$$(s_2, t_2) = (u_2 - u_0, v_2 - v_0)$$

## Computing tangent space basis vectors



object space



texture space

$$\mathbf{d}_1 = \mathbf{p}_1 - \mathbf{p}_0$$

$$(s_1, t_1) = (u_1 - u_0, v_1 - v_0)$$

$$\mathbf{d}_2 = \mathbf{p}_2 - \mathbf{p}_0$$

$$(s_2, t_2) = (u_2 - u_0, v_2 - v_0)$$

$$\mathbf{d}_1 = s_1 \mathbf{P}_u + t_1 \mathbf{P}_v$$

$$\mathbf{d}_2 = s_2 \mathbf{P}_u + t_2 \mathbf{P}_v$$

## Computing tangent space basis vectors

$$\mathbf{d}_1 = s_1 \mathbf{P}_u + t_1 \mathbf{P}_v$$

$$\mathbf{d}_2 = s_2 \mathbf{P}_u + t_2 \mathbf{P}_v$$

## Computing tangent space basis vectors

$$\mathbf{d}_1 = s_1 \mathbf{P}_u + t_1 \mathbf{P}_v$$

$$\mathbf{d}_2 = s_2 \mathbf{P}_u + t_2 \mathbf{P}_v$$

$$\begin{pmatrix} \mathbf{d}_1^T \\ \mathbf{d}_2^T \end{pmatrix} = \begin{pmatrix} s_1 & t_1 \\ s_2 & t_2 \end{pmatrix} \begin{pmatrix} \mathbf{P}_u^T \\ \mathbf{P}_v^T \end{pmatrix}$$

## Computing tangent space basis vectors

$$\mathbf{d}_1 = s_1 \mathbf{P}_u + t_1 \mathbf{P}_v$$

$$\mathbf{d}_2 = s_2 \mathbf{P}_u + t_2 \mathbf{P}_v$$

$$\begin{pmatrix} \mathbf{d}_1^T \\ \mathbf{d}_2^T \end{pmatrix} = \begin{pmatrix} s_1 & t_1 \\ s_2 & t_2 \end{pmatrix} \begin{pmatrix} \mathbf{P}_u^T \\ \mathbf{P}_v^T \end{pmatrix}$$

$$\begin{pmatrix} \mathbf{P}_u^T \\ \mathbf{P}_v^T \end{pmatrix} = \frac{1}{s_1 t_2 - s_2 t_1} \begin{pmatrix} t_2 & -t_1 \\ s_2 & s_1 \end{pmatrix} \begin{pmatrix} \mathbf{d}_1^T \\ \mathbf{d}_2^T \end{pmatrix}$$

We can now normalize these vectors and use them for bump mapping. To compute tangent vectors at a vertex, we average the vectors from the adjacent faces, as we did with normals.

## Normal mapping



original wall



normal map



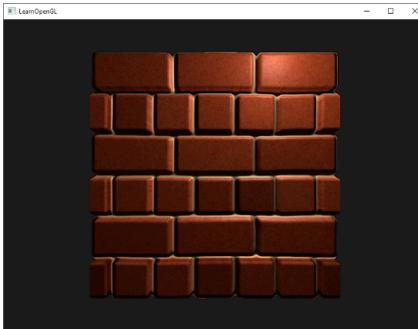
shaded wall

- Store the displaced normals directly. Reduces runtime overhead, at the expense of memory requirements
- (x,y,z) values in the tangent space are stored in the RGB channels. To compute the normal at a fragment, we simply multiply the (interpolated) tangent space basis by (x,y,z)

## Parallax mapping

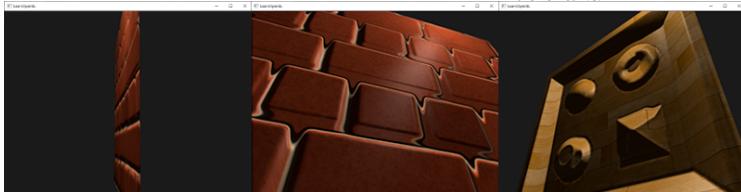


## Parallax mapping

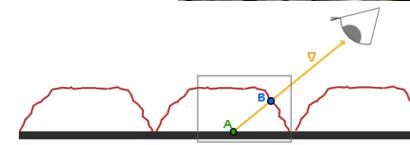


Credit learnopengl.com

The pixel (fragment) hit by ray  $\vec{V}$  returns not its own color, but the color of the texture that would have been returned had the surface followed the red curve.



Can be combined into  
Parallax occlusion mapping



## Relief mapping



normal mapping



relief mapping

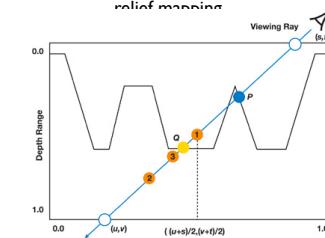
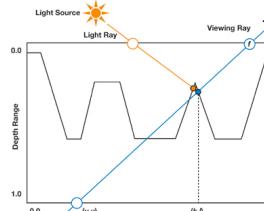
Similar results, different math. Impacts both shading and outcome image.

## Relief mapping



normal mapping

relief mapping



1. Similar results to parallel occlusion map.
2. Different math/algorithm
3. Impacts both shading and outcome image.
4. Credit <https://developer.nvidia.com/gpugems>

## Relief mapping



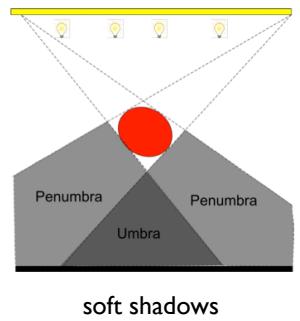
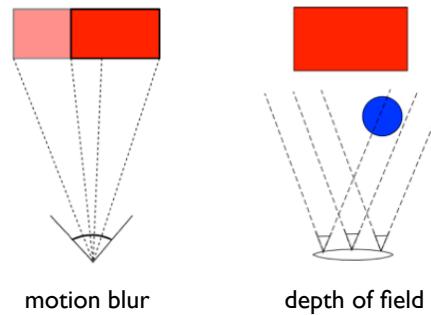
Image from Pollicarpo et al. (2005)

## Other Buffers

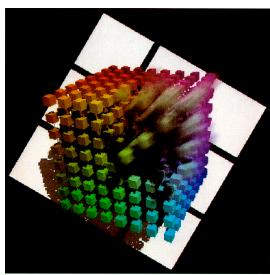
### Accumulation buffer

- High-precision image buffer. Can integrate images that are rendered into the framebuffer. Supports anti-aliasing, motion blur, depth of field, soft shadows, etc.
- 16 bits for each red, green, blue, and alpha component: total of 64 bits per pixel.
- Supports the following operations:
  - Clear: set all values to zero.
  - Add with weight: Each pixel in the drawing buffer is added to the accumulation buffer after being multiplied by a floating-point weight that can be positive or negative.
  - Return with scale: The contents of the accumulation buffer are returned to the drawing buffer after being scaled by a positive floating-point constant
- Can integrate up to 256 images without loss of precision, and even more using weight less than 1.0

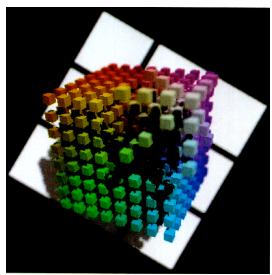
### Accumulation buffer



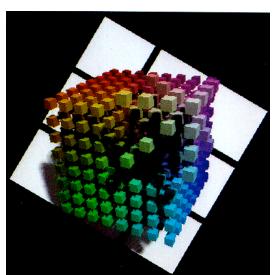
## Accumulation buffer



motion blur



depth of field



soft shadows

Haeblerli and Akeley, SIGGRAPH 1990