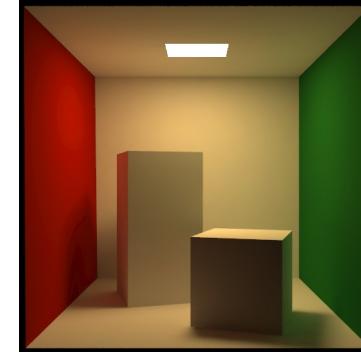


# Rendering Really Large Scenes

Thanks to Siddhartha Chaudhuri and Vladlen Koltun

15 polygons



~500,000 polygons



~3,000,000 polygons



Crysis, 2007

~25,000,000,000 polygons



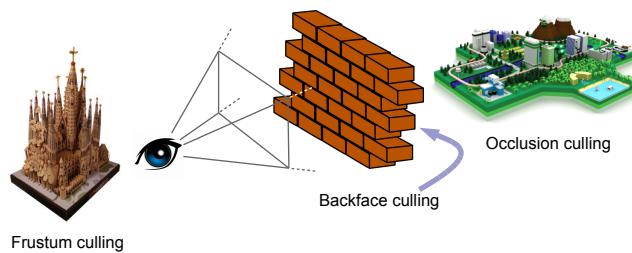
Chaudhuri and Koltun, 2009

## Largescale Rendering Cheat Sheet

- Don't render what you can't see
- Don't render what the display can't resolve
- People won't notice small errors, especially in background objects
- **If all else fails, fog is your best friend :)**

### Don't render what you can't see

- Rasterizing invisible objects is wasteful
- Detect such objects early and ignore (**cull**) them



### Difficulty

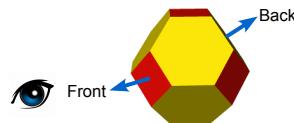
Backface < Frustum << Occlusion

## Backface Culling

- Drop faces on the far side of object meshes
  - Assume face normals consistently point **inside-out**
  - Back faces have **normals pointing away** from the camera
- OpenGL:

```
glEnable(GL_CULL_FACE);
glCullFace(GL_BACK);
```

  - Why would anyone want `glCullFace(GL_FRONT)?`
  - Uses **vertex winding order** to determine front and back faces, not the vertex normals you pass in!
- When does this scheme for back face culling fail?



## Frustum Culling

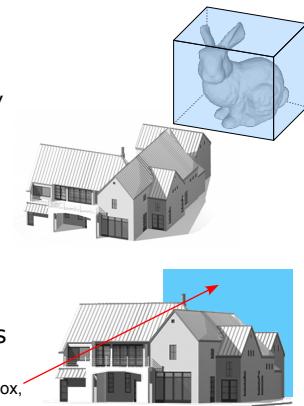
- Test each object against the view frustum
  - Much faster: **test the bounding box** instead
    - If object is visible, no frustum plane can have all 8 corners on invisible side
- Optimization:
  - Group objects hierarchically**
    - Octree* (or *quadtree* for 2.5D scenes)
    - Binary Space Partitioning (BSP) tree* (or a restricted version called a *kd-tree*)
    - Bounding box/sphere hierarchy*
  - Traverse tree top-down and ignore subtrees whose roots fail the bounding box test

## Occlusion Culling

- Whole subject by itself
- General idea:** Ignore background objects covered (**occluded**) by foreground objects
  - (Hardware) Occlusion queries*
  - From-region visibility*
  - Portal-based rendering*

## Hardware Occlusion Queries

- Part of OpenGL/D3D API
- At any time, pretend to draw a dummy shape (say the bounding box of a complex object) and check if any pixels are affected
- Accelerated by hierarchical z-buffer
- Works for dynamic scenes



## From-Region Visibility

- Preprocessing:

- Break scene up into regions
- For each region, compute a **potentially visible set (PVS)** of objects

For gaming, this is performed in a preprocessing stage

- Runtime:

- Detect the region containing the observer
  - Render the objects in the corresponding PVS
- PVS is usually quite conservative, so further culling is needed

## Case Study: Quake

- Preprocessing:

- Level map preprocessed into BSP-tree
- Each leaf node stores potentially visible polygons from that region

- Runtime:

- Leaf node containing player detected by searching the tree (very fast)
- PVS of polygons for this node are rendered
- (BSP-tree is NOT used for back-to-front rendering!) (painter algorithm)

## Portal-Based Rendering

- Suitable for indoor environments

- Divide environment into **cells**, connected by simple polygonal **portals** (doors/windows/...)

- Render:

- Neighboring cells with visible portals (check if projected polygon is within screen limits)
  - Neighbors-of-neighbors with portals visible through the first set of portals
  - ... and so on
- Further culling possible with frusta through portals

## Case Study: Unreal 2



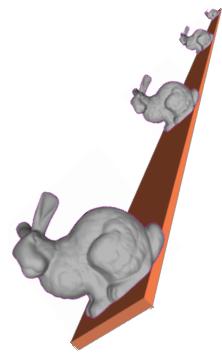
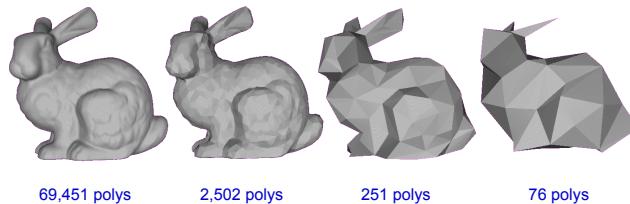
## Don't render what the display can't resolve/people won't notice

- There're only ~1 million pixels on the average screen
- Why spend precious milliseconds rendering the Taj Mahal in exquisite detail if it's only going to take up 10 pixels on the screen?
- **Moral: Simplify distant objects**  
... assuming perspective projection
- **Moral #2:** Since people are usually not too interested in the background, you can **simplify over-aggressively**

For every object, choose the **simplest possible representation** that will look nearly the same as the original.

## Levels of Detail (LOD)

- **Coarser representations for distant objects**
  - Hierarchy of representations of the same object at different resolutions
  - The same idea can also be used for textures (**mipmapping**)



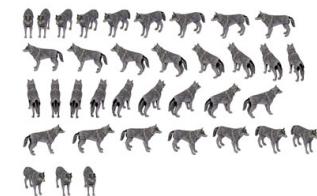
## Environment Maps

- Very distant stuff looks the same from anywhere within reasonable limits
- Pre-render distant objects (including the sky) out to a 360° image
- Texture-map it onto a bounding cube at runtime



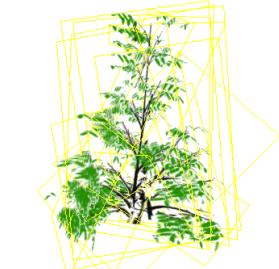
## Image-Based Rendering

- Render complex objects to images and texture-map them to simple proxy shapes (**impostors**)
  - Environment mapping is a specific example
- **Billboards/sprites:** Textured quads always facing the viewer
  - Single image is valid if viewer doesn't move much

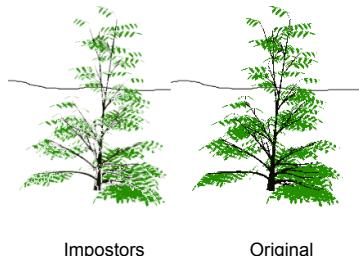


## Image-Based Rendering

Image base rendering - use real images in computer-generated scene



Tree decomposed into a cloud of texture-mapped planar slices

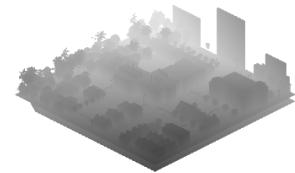


Décoret, Sillion, Durand and Dorsey 2002

- We create several (partially overlapping) images of portions of the tree.
- Each captures more or less a plane in the tree
- During rendering time, we decide which of the image planes to use (more or less)
- Note the use of transparent pixels  $\alpha = 0$

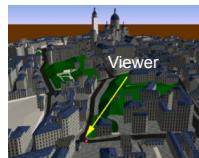
## Adding Depth to Images

- Store the **depth map** as well as the color
- Impostor is **heightfield** defined by the depth map
- Fixes parallax errors (impostor is still valid when viewing position changes significantly)
- What are the drawbacks?

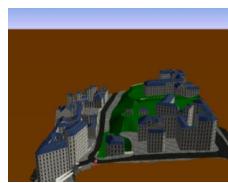


Think about maintaining the image as viewer moves

## Images + Geometry



=



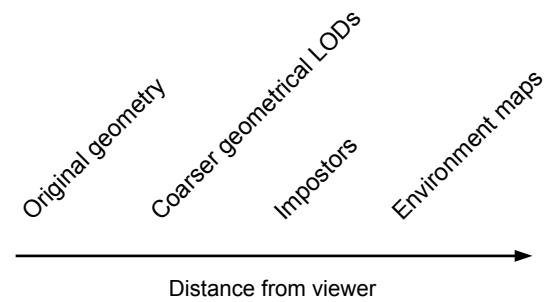
+



Foreground

Background

## LOD spectrum (not exhaustive or exact!)



## Speed Gains (back of the envelope)

- Backface culling:  $\sim 2x$
- Frustum culling:  $\sim 5x$  (varies inversely with FOV)
- Occlusion culling: can be huge, but no guarantees at all except in special cases, e.g. portal-based indoor environments
- Levels-of-detail: a uniformly dense scene of radius  $r$  takes  $O(\log r)$  time to render (prove!)

Lesson of the day: invest in good LODs!!!