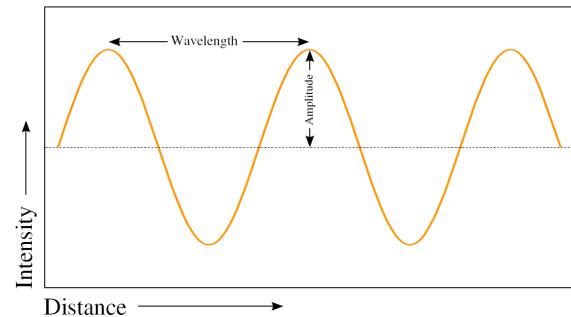
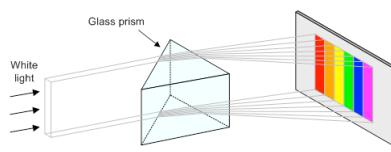
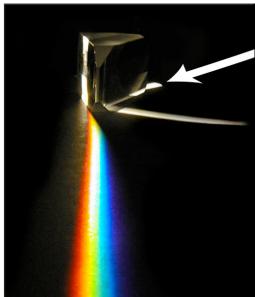


CSC 433/533 Computer Graphics Colors, Sampling, Aliasing, Convolutions

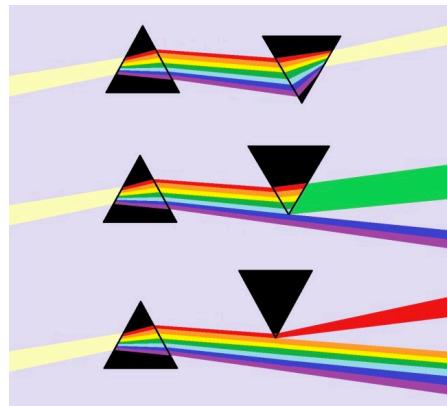
Light



Isaac Newton, 1666



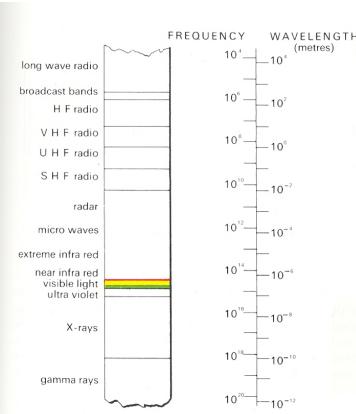
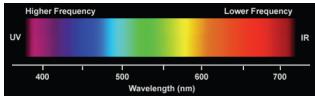
<http://www.webexhibits.org/colorart/bh.html>
<https://www.clivemaxfield.com/diycalculator/popup-m-cvision.shtml>



<http://www.thestargarden.co.uk/Newton's-theory-of-light.html>

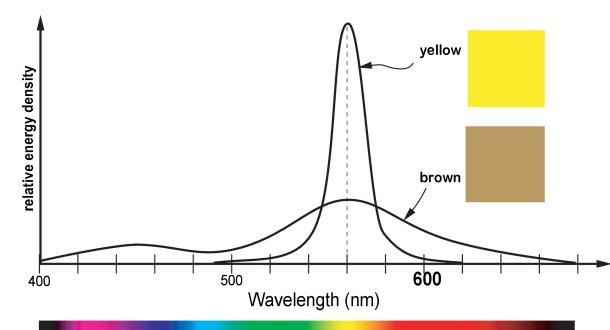
Recall: Light is Electromagnetic Radiation

- Visible spectrum is “tiny”
- Wavelength range: 380-740 nm



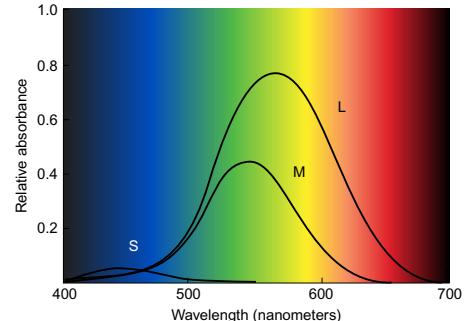
Recall: Color != Wavelength

- But rather, an integral over the wavelengths of the energy encoded of some **power spectrum**



Color and Perception

Recall: We have three types of cones (Short, Medium, and Long)



Colin Ware, Information Visualization: Perception for Design

Hunters

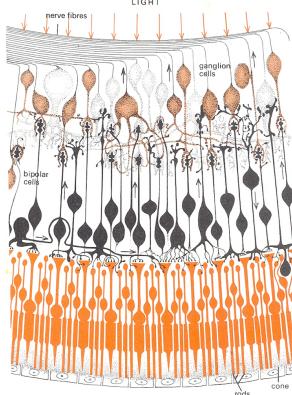


Gatherers



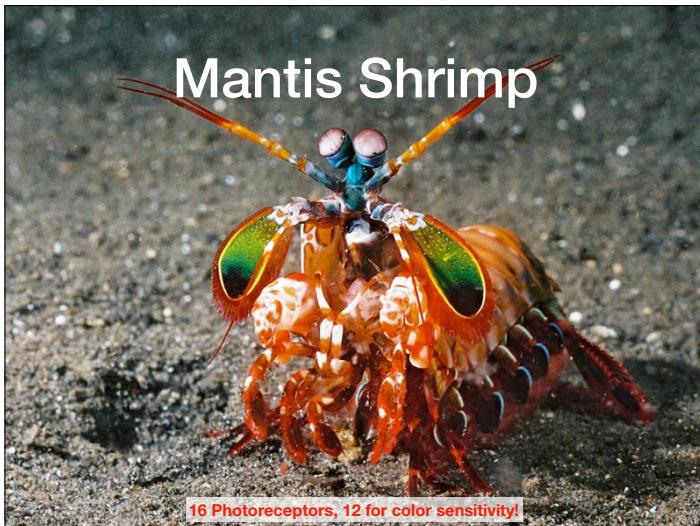
Photoreceptors

- Rods** (detect low-light / scotopic vision)
 - Approximately 100-150 million rods (Non-uniformly distributed across the retina)
 - Sensitive to low-light levels (scotopic vision)
- Cones** (detect day-light / photopic vision)
 - Approximately 6-7 million cones.
 - Detect color with 3 different kinds:
 - Red (L cone) : 564-580nm wavelengths (65% of all cones)
 - Green (M cone) : 534-545nm (30% of all cones)
 - Blue (S cone) : 420-440nm (5% of all cones)



Trichromacy

- Our 3 cones cover the visible spectrum (theoretically, all we might have are 2 though)
- Most birds, some fish, reptiles, and insects have 4, some as many as 12 (e.g. the mantis shrimp)
- This is a “reason” why many of our acquisition devices and displays use 3 channels, and why many of our color spaces are three dimensional



Key Idea: Perception of color



Ultimately, color is a perceptual phenomenon, we all perceive it differently

Color Models

Color Terminology

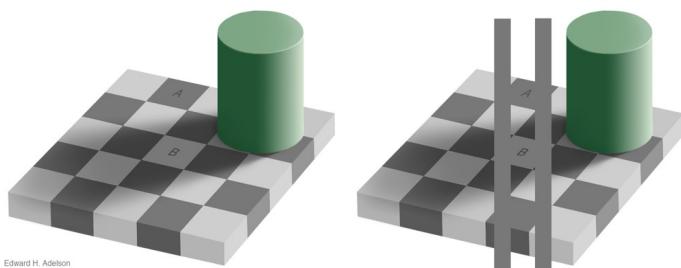
- **Color Model**

- Is an abstract mathematical system for representing color.
- Is often 3-dimensional, but not necessarily.
- Is typically limited in the range of colors they can represent and hence often can't represent all colors in the visible spectrum

- **Gamut or Color Space**

- The range of colors that are covered by a color model.

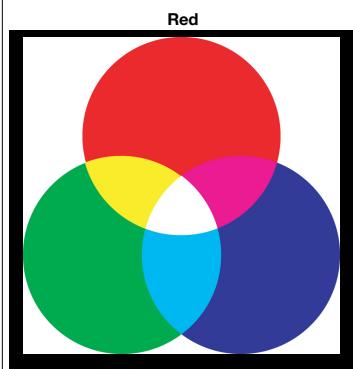
Simultaneous Contrast



Edward H. Adelson

http://persci.mit.edu/_media/gallery/checkershadow_double_full.jpg

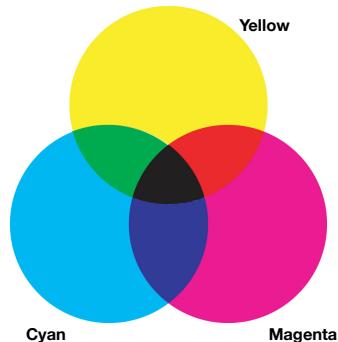
Light Mixing



- **Additive** mix of colored lights (start with black)
 - Add up wavelengths of light to make new colors
- Primary: RGB
- Secondary: CMY (cyan, magenta yellow)
- Neutral = R + G + B
- Commonly used by monitors, projectors, etc.

Ink Mixing

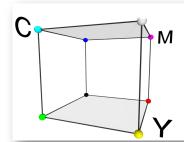
Different game, since we start with white page rather than a black screen
Each color filters the light that is reflected from the white page.



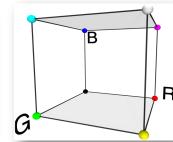
- Subtractive mix of transparent inks
 - Start with white and other wavelengths are selectively filtered.
 - The Yellow region does not completely prevent reflection of light from the white page. But it TENDS (depending on transparency) to filter others frequencies)
- Primary: CMY (Cyan, Magenta, Yellow)
- Secondary: RGB
- ~Black: C + M + Y
- In practice, we use CMYK, with some amount K of black ink, to get true black

Converting from RGB to CMY

- Assuming RGB values are normalized (all channels between [0,1]), the exact same color in CMY space can be found by inverting:

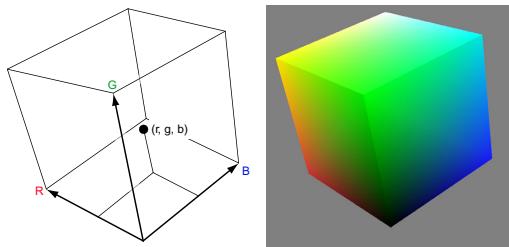


$$\begin{bmatrix} C \\ M \\ Y \end{bmatrix} = \begin{bmatrix} 1-R \\ 1-G \\ 1-B \end{bmatrix}$$



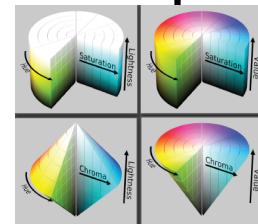
RGB Color Space

- Additive, useful for computer monitors
- Not perceptually uniform
 - For example, more “greens” than “yellows”



HSL, HSV Color Space

- Hue - what people think of as color (color, normalized by sensitivity). Specified as an angle between 0° .. 360°
- Saturation - purity, distance from grey
- Brightness - is proportional to the gap $\text{Max}(R, G, B) - \text{Min}(R, G, B)$
 - Also called Chroma
- Lightness - from dark to light (how many photons, alternatively, add more sources of light)
 - Also Brightness or Value or Intensity



Hue wheel (credit: Wiki)
(not a single frequency)



Originally used for TV broadcasting: The broadcasting signal should be used by both B&W TV and by color TV. The latter ones just see the L of the HSL.

HSL was invented for television in 1938 by Georges Valensi television encoding including NTSC, PAL and SECAM and all major digital broadcast systems and is the basis for composite video.

Conversion from RGB to HSB

- Assuming RGB values are normalized (all channels between [0,1]), the exact same color in HSB space can be found by first figuring out which channel (R,G, or B) has the max intensity

$$H = \begin{cases} \text{undefined} & \text{if } \max = \min, \\ 60 \times \frac{G-B}{\max-\min} & \text{if } \max = R \text{ and } G \geq B, \\ 60 \times \frac{G-B}{\max-\min} + 360 & \text{if } \max = R \text{ and } G < B, \\ 60 \times \frac{B-R}{\max-\min} + 120 & \text{if } \max = G, \\ 60 \times \frac{R-G}{\max-\min} + 240 & \text{if } \max = B. \end{cases}$$

(Note: this method returns H as a value between 0° and 360°)

$$S = \begin{cases} 0 & \text{if } \max = 0, \\ 1 - \frac{\min}{\max} & \text{otherwise} \end{cases}$$

$B = \max // 'B'$ for “brightness”. Not ‘B’ for “blue”

Anti-Aliasing and Signal Processing

Sampling, Smoothing and Convolutions

Recall: Images are Functions

Domains and Ranges

- All functions have two components, the **domain** and **range**. For the case of images, $I: R \rightarrow V$
- The domain is:
 - R , is some rectangular area ($R \subseteq \mathbb{R}^2$)
- The range is:
 - A set of possible values.
 - ...in the space of color values we're encoding

Concept for the Day: Pixels are Samples of Image Functions

Image Samples

- Each pixel is a sample of what?
- One interpretation: a pixel represents the intensity of light at a single (infinitely small point in space)
- The sample is displayed in such a way as to spread the point out across some spatial area (drawing a square of color)

Continuous vs. Discrete

- Key Idea: An image represents data in either (both?) of
 - Continuous domain: where light intensity is defined at every (infinitesimally small) point in some projection
 - Discrete domain, where intensity is defined only at a discretely sampled set of points.
- This seem like a philosophical discussions without clear practical applications. Surprisingly, it has very concrete algorithmic applications.

Converting Between Image Domains

- When an image is acquired, an image is **sampled** from some continuous domain to a discrete domain.
- **Reconstruction** converts digital back to continuous.
- The reconstructed image can then be **resampled** and **quantized** back to the discrete domain.

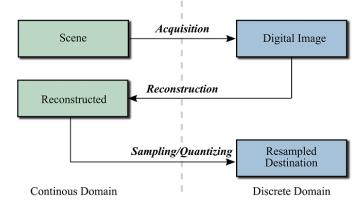


Figure 7.7. Resampling.

```

//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//copy the pixels over
for (let row = 0, row < H; row++) {
    for (let col = 0; col < W; col++) {
        let index = row*W + col;
        let index2 = (k*row)*W + (k*col);
        output[index2] = input[index];
    }
}

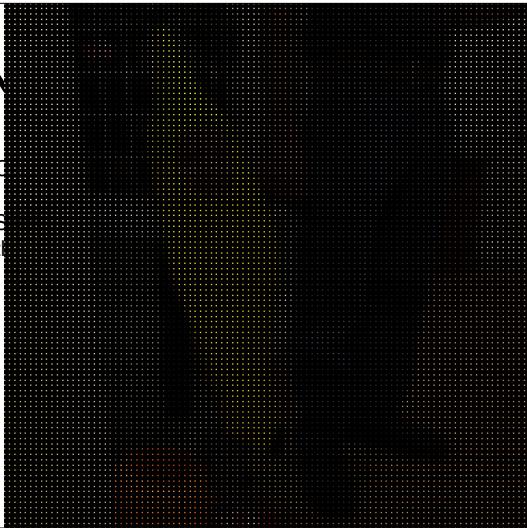
```

Naive Image Rescaling Code



Naive Image Rescaling

- Good for learning
- Simple to understand



What's the Problem?

- The output image has gaps!
- Why: we skip a many of the pixels in the output.
- Why don't we fix this by changing the code to at least put some color at each pixel of the output?

```

//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//copy the pixels over
for (let row = 0, row < H; row++) {
    for (let col = 0; col < W; col++) {
        let index = row*W + col;
        let index2 = (k*row)*W + (k*col);
        output[index2] = input[index];
    }
}

```

Naive Image Rescaling Code

```

//scale factor
let k = 4;

//create an output greyscale image that is both
//k times as wide and k times as tall
Uint8Array output = new Uint8Array((k*W)*(k*H));

//Loop over each output pixel instead.
for (let row = 0, row < k*H; row++) {
    for (let col = 0; col < k*W; col++) {
        let index = (row/k)*W + (col/k);
        let index2 = row*k*W + col;
        output[index2] = input[index];
    }
}

```

“Inverse” Image Rescaling Code

Red arrows highlight the changes made to the original code:

- A red arrow points to the line `for (let row = 0, row < k*H; row++) {`, indicating the loop range was changed from `row < H` to `row < k*H`.
- A red arrow points to the line `let index = (row/k)*W + (col/k);`, indicating the calculation of the output index was modified to account for the scaling factor `k`.

Inverse Image Rescaling



400x400 image

Not great, but could become worse



100x100 image

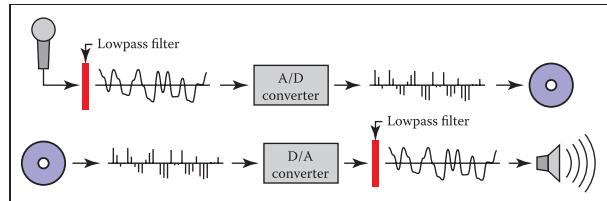
What's the Problem?

- The output image is too “blocky”
- Why: because our image reconstruction rounds the index to the nearest integer pixel coordinates
 - Rounding to the “nearest” is why this type of interpolation is called **nearest neighbor interpolation**

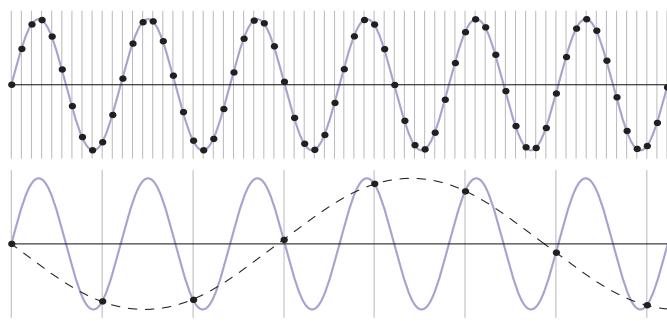
Sampling Artifacts / Aliasing

Motivation: Digital Audio

- Acquisition of images takes a continuous object and converts this signal to something digital
- Two types of artifacts:
 - **Undersampling** artifacts: on acquisition side
 - **Reconstruction** artifacts: when the samples are interpreted



Undersampling Artifacts



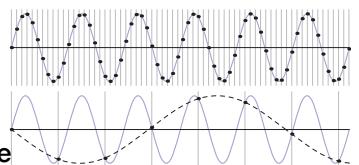
<http://youtu.be/0k2lhYk6Lfs?rel=0>



Handwaved

Shannon-Nyquist Theorem

(not needed for the exam)

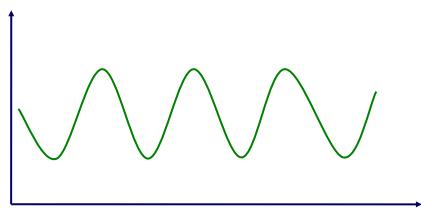


- The sampling frequency must be double the highest frequency of the content.
- If there are any higher frequencies in the data, or the sampling rate is too low, **aliasing**, happens
- Named this because the discrete signal “pretends” to be something lower frequency

S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

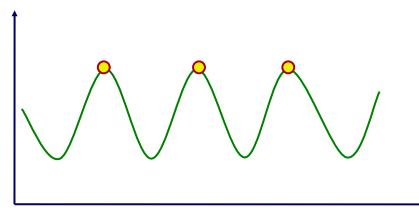
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

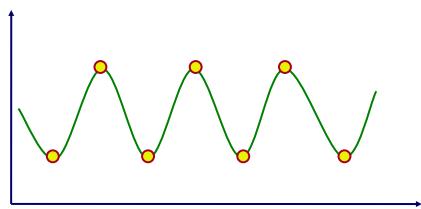
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

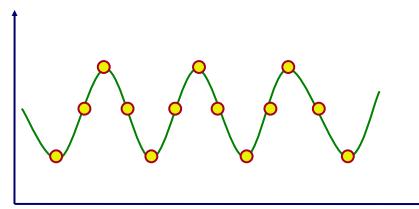
- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?



S-N Theorem Illustrated

How many samples are enough to avoid aliasing?

- How many samples are required to represent a given signal without loss of information?
- What signals can be reconstructed without loss for a given sampling rate?

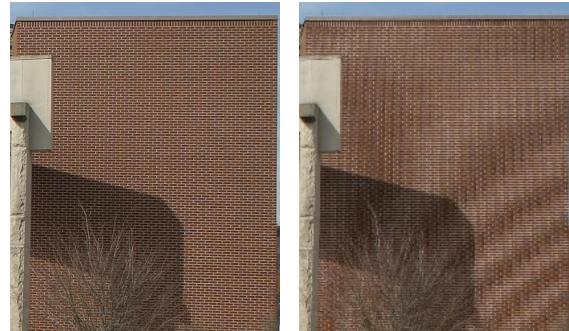


Aliasing in images

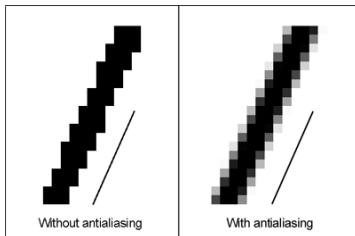
Two outcomes of under-sampling

- 1) Moire Pattern
- 2) Rasterization

Moire Patterns



Aliasing for edges



Each pixel is effected by nearby pixels
For example, even though the input image image is black/white,
We allow grey values for output pixels.

Convolution

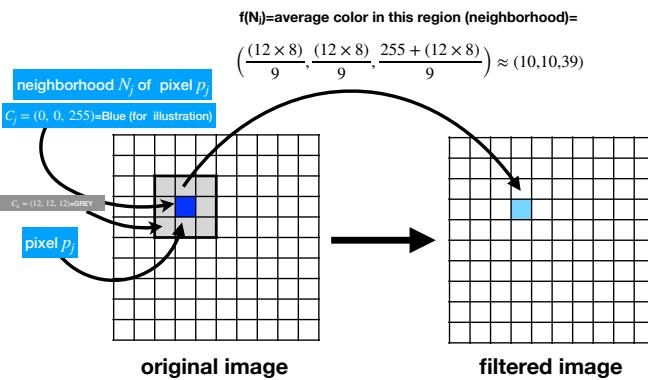
Convolutions will be used for several applications in this course :

- Anti-Aliasing,
- Sharpening of images
- Dynamic Range (overcome limitations of monitor: We only have 2^8 levels of intensity, which is very far from sufficient - but we will study how to display images)

Big idea. Image could be improved, if when setting the value of a pixel, we also take into account values of nearby pixels



Neighborhood Filtering (Schematic)



An Example: Mean Filtering

• Mean filters sum all of the pixels in a local neighborhood N_i and divide by the total number, computing the average pixel.

• Said another way, we replace each pixel as a linear combination of its neighbors (with equal weights)

• To find the new color of a pixel p_j , we will look at N_j , defined as the (say) 3×3 neighborhood of the pixel p_j , and set

• Where the N_i is a square, we call these **box filters**

• Think about it as a weighted average:

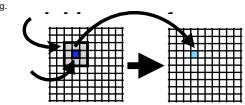
$$f(N_i) = \sum_{\text{pixels } p_k \text{ in the region } N_i} w_k C_k = \sum_{\text{pixels } p_k \text{ in the region } N_i} \frac{1}{9} C_k$$

• The weights $w_1 \dots w_9$ are convex combination. Meaning that they are all positive, and $w_1 + w_2 + \dots + w_9 = 1$. For example, $w_1 = w_2 = w_3 = \frac{1}{3}$ (convex combination)

• Remember: The input matrix and the output matrix have the same size (in this case). This is **not** rescaling.

• Refer to the geogebra app <https://www.geogebra.org/m/cetpwaw>

• The term filter is very common, but might be very confusing. We don't necessarily filter out anything.



Box Filtering



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$

The matrix of weights is called a **Kernel**



$$1/9 * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Box Filtering



$$1/9 * \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

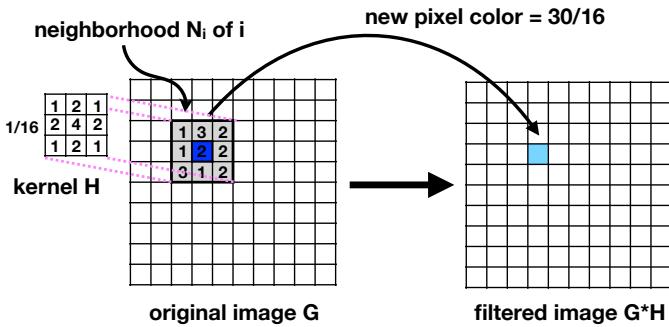


$$1/25 * \begin{matrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{matrix}$$



Convolution

- This process of adding up pixels multiplied by various weights is called **convolution**. We denote the result by (confusion warning) the symbol * See example below.



Kernels

- Convolution employs a rectangular grid of coefficients, (that is, weights) known as a **kernel**
- Kernels are like a neighborhood mask, they specify which elements of the image are in the neighborhood and their relative weights.
- A kernel is a set of weights that is applied to corresponding input samples that are summed to produce the output sample.
- For **smoothing** purposes, the sum of weights must be 1 (convex combination)

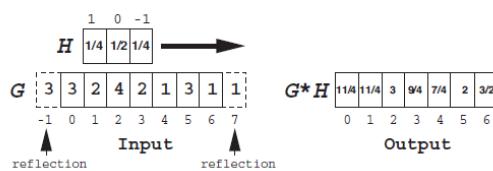
$$\frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & \textcolor{red}{1} & 1 \\ 1 & 1 & 1 \end{pmatrix} \quad \frac{1}{13} \begin{pmatrix} 1 & 1 & 1 \\ 1 & \textcolor{red}{5} & 1 \\ 1 & 1 & 1 \end{pmatrix} \quad \frac{1}{37} \begin{pmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 2 & 2 & 2 & 1 \\ 1 & \textcolor{red}{2} & \textcolor{red}{5} & 2 & 1 \\ 1 & 2 & 2 & 2 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$

One-dimensional Convolution

- Can be expressed by the following equation, which takes a filter H and convolves it with G :

$$\hat{G}[i] = (G * H)[i] = \sum_{j=i-n}^{i+n} G[j]H[i-j], \quad i \in [0, N-1]$$

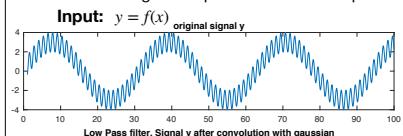
- Equivalent to sliding a window



Low-pass and high-pass filtering

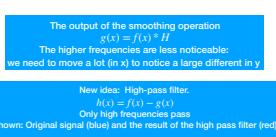
The smoothing operation is always a low pass filter.

Only lower frequencies could pass.
It removes higher frequencies from the input.



We convolved the original signal $f(x)$ a smoothing kernel H . For example

$$g(x) = \frac{f(x-1) + f(x) + f(x+1)}{3}$$

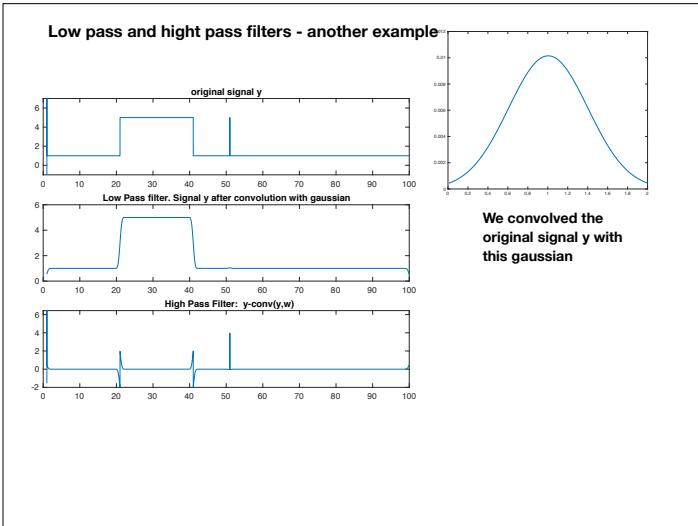


New idea: High-pass filter.
 $h(x) = f(x) - g(x)$
Only high frequencies pass
(shown: Original signal (blue) and the result of the high pass filter (red))

We remove (subtract) from the signal all lower frequencies

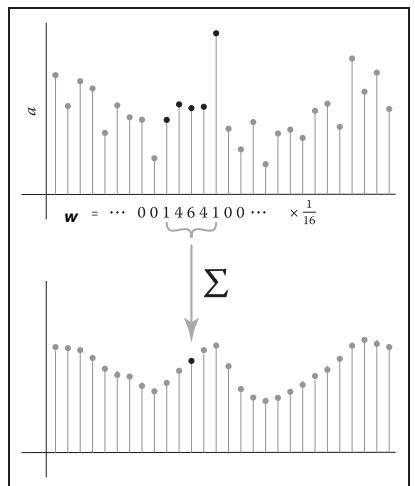
Twitter - could move very fast, but only small distances





Convolution is a Moving, Weighted Average

- Getting used to the new notation:
 $b[i] = \frac{1}{3}(a[i-1] + a[i] + a[i+1])$ Vi
 is similar to writing $b = a * w$, where
 $b[i] = (a * w)[i] = \sum_{j=1}^{m+k} a[i-j+2] \cdot w[j]$ and
 $w[1]=w[2]=w[3]=1/3$
 - Commonly $(a * w)[i] = \sum_{j=i-k}^{i+k} a[j]w[i-j]$
 - For example, $w[-1]=w[0]=w[1]=1/3$
 - Note that we did not define exactly what are the first and last values



2-Dimensional Version

- Given an image a and a kernel b with $(2r+1)^2$ values, the convolution of a with b is given below as $a * b$:

$$(a \star b)[i, j] = \sum_{i'=i-r}^{i+r} \sum_{j'=j-r}^{j+r} a[i', j'] b[i - i', j - j']$$

- The $(i-i')$ and $(j-j')$ terms can be understood as reflections of the kernel about the central vertical and horizontal axes.
 - The kernel weights are multiplied by the corresponding image samples and then summed together.

A Note on Indexing

- Convolution **reflects** the filter to preserve orientation.
 - Correlation does **not** have this reflection.
 - But we often use them interchangeably since most kernels are symmetric!!

**Convolution reflects
and shifts the kernel**

Given kernel H =

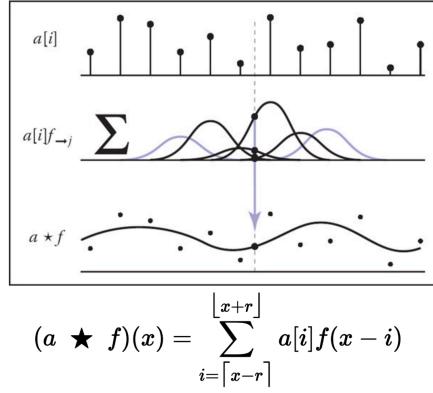
1	2	3
4	5	6
7	8	9

0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 1 2 3 0
 0 0 0 0 0 0 0 0 0 0 4 5 6 0
 0 0 0 1 2 3 0 0 0 0 7 8 9 0
 0 0 0 4 5 6 0 0 0 0 0 0 0 0
 0 0 0 7 8 9 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0

G*H

Convolution Can Also Convert from Discrete to Continuous

- Discrete signal a
 - Continuous filter f
 - Output a^*f defined on positions x as opposed to discrete pixels i



Filtering helps to reconstruct the signal better when rescaling



Inverse Rescaling



Reconstructed w/ Discrete-to-Continuous

Types of Filters: Smoothing

Smoothing Spatial Filters

- Any weighted filter with positive values will smooth in some way, examples:

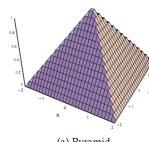
$$\frac{1}{9} \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$\frac{1}{16} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

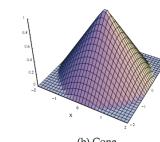
- Normally, we use integers in the filter, and then divide by the sum (computationally more efficient)
- These are also called **blurring** or **low-pass filters**

Smoothing Kernels

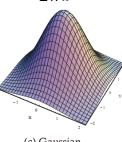
$$f(x, y) = -\alpha \cdot \max(|x|, |y|)$$



(a) Pyramid.



(b) Cone.



(c) Gaussian.

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

$$f(x, y) = -\alpha \cdot \sqrt{x^2 + y^2}$$

1	2	3	2	1
2	4	6	4	2
3	6	9	6	3
2	4	6	4	2
1	2	3	2	1

(a) Pyramid.

0	0	1	0	0
0	2	2	2	0
1	2	5	2	1
0	2	2	2	0
0	0	1	0	0

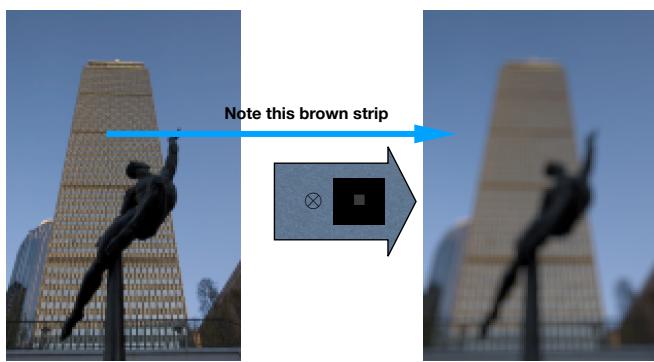
(b) Cone.

1	4	7	4	1
4	16	28	16	4
7	28	49	28	7
4	16	28	16	4
1	4	7	4	1

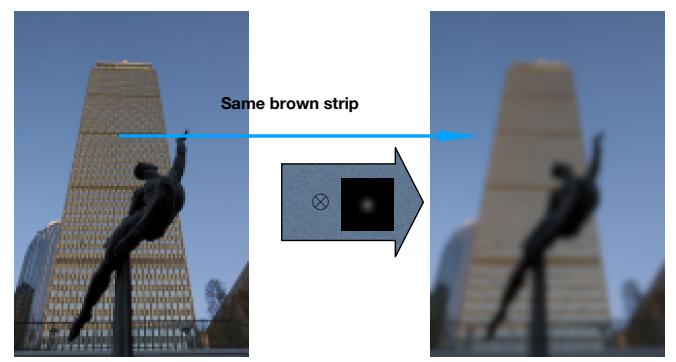
(c) Gaussian.

Table 6.1. Discretized kernels.

Box Filter



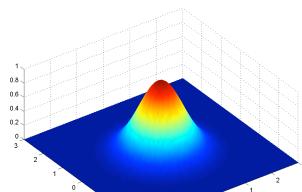
Gaussian Filter



Gaussians

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- Gaussian kernel is parameterized on the standard deviation σ
- Large σ 's reduce the center peak and spread the information across a larger area
- Smaller σ 's create a thinner and taller peak
- Gaussians are smooth everywhere.
- Gaussians have infinite **support**
 - >0 everywhere
- But often truncate to 2σ or 3σ
- Volume =1 (sum of weights =1)



http://en.wikipedia.org/wiki/Gaussian_function

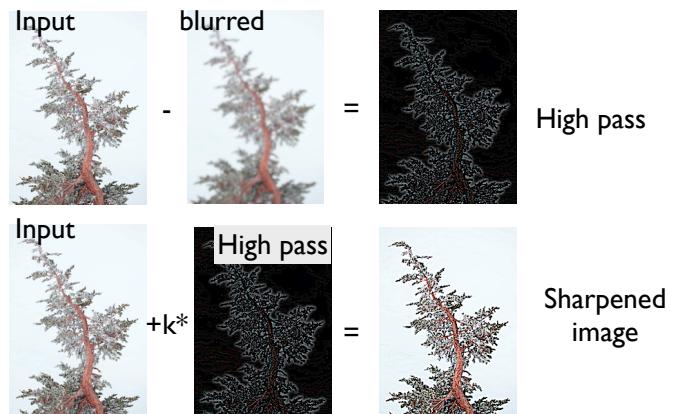
Smoothing Comparison



Figure 6.10. Smoothing examples.

Types of Filters: Sharpening

Sharpening (Idea)



Another example

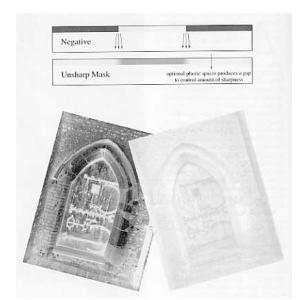
Original Image, Imaged convolved



Left: difference (only boundaries are non-black)
Right: Imaged minus differences convolved

Unsharp Masks

- Sharpening is often called “unsharp mask” because photographers used to sandwich a negative with a blurry positive film in order to sharpen



<http://www.tech-diy.com/UnsharpMasks.htm>

Edge Enhancement

- The parameter α controls how much of the source image is passed through to the sharpened image.

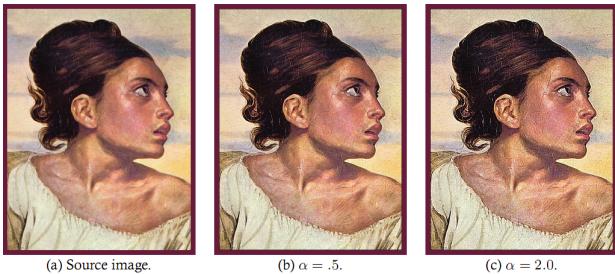


Figure 6.20. Image sharpening.

Defining Edges

- Sharpening uses negative weights to enhance regions where the image is changing rapidly
 - These rapid transitions between light and dark regions are called **edges**
- Smoothing reduces the strength of edges, sharpening strengthens them.
 - Also called **high-pass filters**
- Idea: smoothing filters are weighted averages, or integrals. Sharpening filters are weighted differences, or derivatives!

Edges

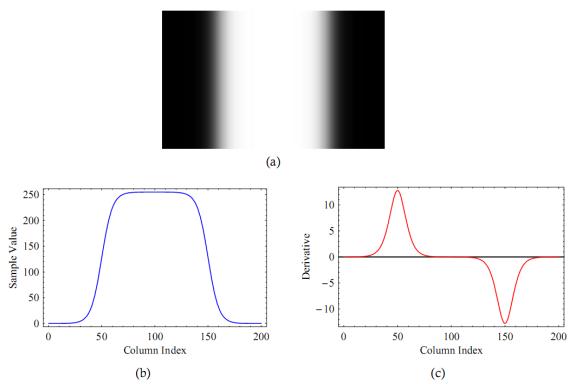
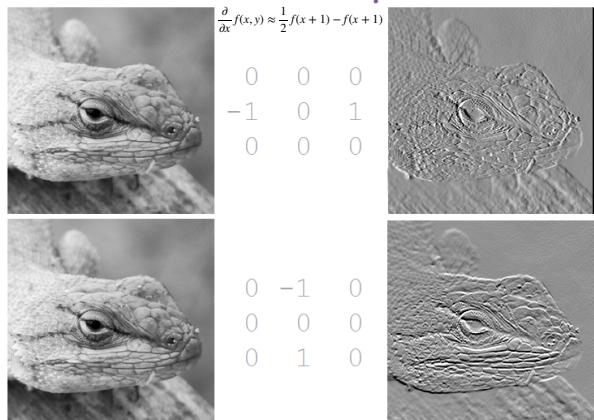


Figure 6.11. (a) A grayscale image with two edges, (b) row profile, and (c) first derivative.

Taking Derivatives with Convolution (just in case you studied calculus. Not required)



Gradients with Finite Differences (just in case you studied calculus. Not required)

- These partial derivatives approximate the image gradient, ∇I .
- Gradients are the unique direction where the image is changing the most rapidly, like a slope in high dimensions
- We can separate them into components kernels G_x , G_y . $\nabla I = (G_x, G_y)$

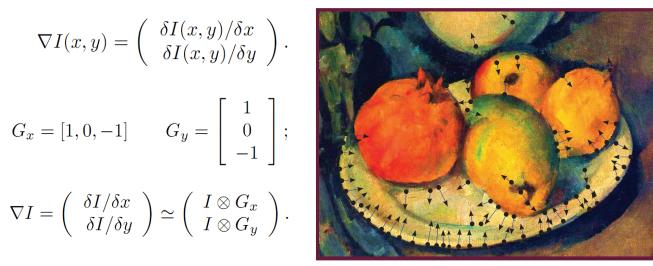


Figure 6.12. Image gradient (partial).

128	187	210	238	251
76	121	193	225	219
66	91	110	165	205
47	81	83	119	157
41	59	63	75	125

(a) Source Image.

117	104	26	
44	74	95	
36	38	74	

(b) $\delta I/\delta x$.

-96	-100	-73	
-40	-110	-106	
-32	-47	-90	

(c) $\delta I/\delta y$.

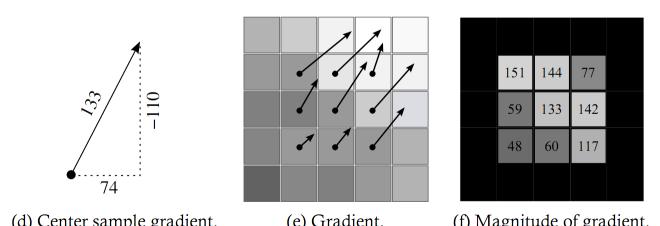
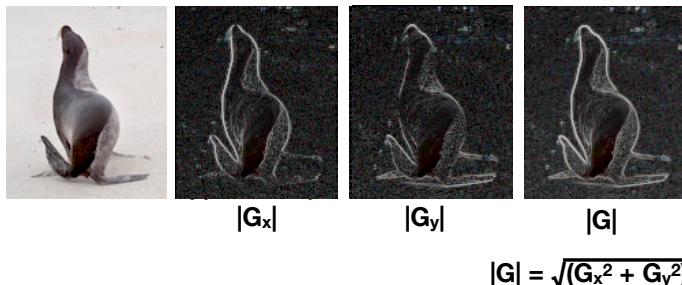


Figure 6.14. Numeric example of an image gradient.

Gradients G_x, G_y



Second Derivatives (Sharpening, almost)

- Partial derivatives in x and y lead to two kernels:

$$\frac{\partial^2 f}{\partial x^2} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

and, similarly, in the y-direction we have

$$\frac{\partial^2 f}{\partial y^2} = f(x, y+1) + f(x, y-1) - 2f(x, y)$$

0	1	0
1	-4	1
0	1	0

1	1	1
1	-8	1
1	1	1

Compare with
Sharpening filter:
unbalanced counts!

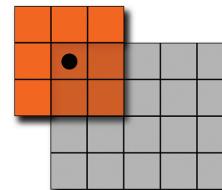
$$\begin{bmatrix} -\alpha & -\alpha & -\alpha \\ -\alpha & (9+8\alpha) & -\alpha \\ -\alpha & -\alpha & -\alpha \end{bmatrix}$$

1	1	1
1	-9	1
1	1	1

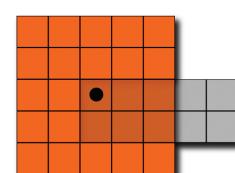
Boundaries

Handling Image Boundaries

- What should be done if the kernel falls off the boundary of the source image as shown in the illustrations below?



(a) Kernel at $I(0, 0)$.



(b) Kernel larger than the source.

Figure 6.4. Illustration of the edge handling problem.

Handling Image Boundaries

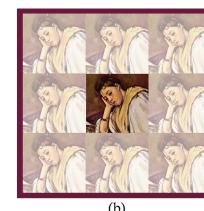
- When pixels are near the edge of the image, neighborhoods become tricky to define
- Choices:
 - Shrink the output image (ignore pixels near the boundary)
 - Expanding the input image (padding to create values near the boundary which are “meaningful”)
 - Shrink the kernel (skip values that are outside the boundary, and reweigh accordingly)

Boundary Padding

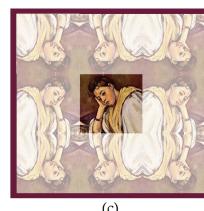
- When one pads, they pretend the image is large and either produce a constant (e.g. zero), or use circular / reflected indexing to tile the image:



(a)



(b)



(c)

Figure 6.5. (a) Zero padding, (b) circular indexing, and (c) reflected indexing.